

Assignment 1

Problem 1 (15 points) In the MAX-COVERAGE problem, we are given m sets S_1, \dots, S_m , each a subset of a universe U . We are also given an integer parameter $k \geq 1$. The objective is to pick k out of these m sets, indexed by I with $|I| = k$ such that $\bigcup_{j \in I} S_j$ is maximized. Describe a natural greedy algorithm for this problem and prove that it is an $(1 - 1/e)$ -approximation.

Given $U = \{e_1, e_2, \dots, e_n\}$
 $F = \{S_1, S_2, \dots, S_m\}$ and $k \geq 1$

Greedy Algo

greedy-Max-Coverage(U, F, k):

$$X \leftarrow U$$

$$I \leftarrow \emptyset$$

while $k > 0$:

pick S_j^* s.t. $|S_j^* \cap X|$ is max

$$X \leftarrow X - S_j^*$$

$$I \leftarrow I \cup \{S_j^*\}$$

$$k = k - 1$$

return I

To Prove: $(1 - 1/e) \text{OPT} \leq T$

Notations: $\text{OPT} = \text{no. of elements covered by the optimal solution}$

$T = \text{no. of elements covered by greedy}$

$n_j^* = \text{no. of elements covered in the } j^{\text{th}} \text{ iteration}$

$x_j^* = \text{no. of elements remaining to be covered at the beginning of the } j^{\text{th}} \text{ iteration}$

$z_j^* = \text{no. of elements covered after } j \text{ rounds}$

$$x_{j+1}^o = x_j^o - n_j^o \quad \text{and} \quad n_j^o \geq \frac{x_j^o}{k} \geq \frac{\text{OPT} - \sum_{i=0}^{j-1} n_i^o}{k} \quad \text{--- (1)}$$

[no. of elements left to be covered from $U \geq$
 no. of _____ from OPT
 since $\text{OPT} \leq |U|$]

$$\text{Also, } n_j^o \leq n \left(1 - \frac{1}{k}\right)^{j-1} \quad \text{--- (2) (from SET-COVER)}$$

$$Z_j^o = \overset{n}{\underset{|U|}{\circlearrowleft}} - x_{j+1}^o$$

$$= n - x_{j+1}^o$$

$$\leq n - n \left(1 - \frac{1}{k}\right)^j = n \left(1 - \left(1 - \frac{1}{k}\right)^j\right) \quad \text{(using (2))} \quad \text{--- (3)}$$

$$\text{Also, } Z_j^o = \sum_{i=1}^j n_i^o$$

$$\geq \sum_{i=1}^j \frac{\text{OPT} - Z_{i-1}^o}{k}$$

$$\geq \sum_{i=1}^j \frac{\text{OPT} - n \left[1 - \left(1 - \frac{1}{k}\right)^{i-1}\right]}{k} \quad \text{(using (3))}$$

$$= \frac{1}{k} \text{OPT} - \frac{n}{k} \left[j - \sum_{i=1}^j \left(1 - \frac{1}{k}\right)^{i-1}\right]$$

Replacing j with k ,

$$Z_k \geq \text{OPT} - \frac{n}{k} \left[k - \left[\frac{1 - (1 - \frac{1}{k})^k}{1 - (1 - \frac{1}{k})}\right]\right] \quad \text{(sum of GP)}$$

$$\geq \text{OPT} - \frac{n}{k} \left[k + k(1 - 1/e) \right]$$

$$\geq \text{OPT} - \underbrace{n}_{\geq \text{OPT}} \underbrace{\left[2 - 1/e \right]}_{\geq 1/e}$$

$$\geq \text{OPT} - \text{OPT} \approx 1/e \geq (1 - 1/e) \text{OPT}$$

$$\Rightarrow \boxed{Z_k (=T) \geq (1 - 1/e) \text{OPT}}$$

Problem 2 (15 points) We consider the makespan minimization problem with an additional constraint. The basic setting is the same - you have a set of n jobs with processing times $p_j = 1$ for every job j (yes, this is an easier setting with all jobs of unit size) and m identical parallel machines. In addition, now you are given *precedence constraints* among jobs - $j > j'$ means that j' can be processed only after processing j . Note that precedence constraints are transitive. We can think of the precedence constraints as inducing a collection of directed acyclic graphs on the jobs. The task is to find a schedule of the jobs to minimize makespan. Note here that now it is not enough to just assign a job to a machine (like we did for basic makespan in class), rather you need to specify exact time-points where to process the jobs (you are not allowed to break jobs).

Extend the basic greedy algorithm to give a 2-approximation for this problem.

2. First, we perform some preprocessing as follows:-

1. Create a DAG based on the precedence constraints
 - each node represents a set of jobs (will be used later) and initially has only 1 job
 - each node has a weight = sum of p_j 's in the node
 - each node has a min. starting time attribute (initially 0 for all nodes with no incoming edges and ∞ for the rest)

2. Compress the DAG

- if there are 2 nodes which have only 1 incoming edge and 1 outgoing edge, AND have an edge between them THEN combine them

New weight = sum of the weights of two nodes

New MST = min of MST of the 2 nodes

- repeat till no nodes can be combined

e.g.



Greedy Algorithm:

1. Let the min MST = t & nodes in DAG
2. Sort all nodes with MST = t in decreasing order of wt
3. Schedule them based on List Scheduling
 - while scheduling make sure the start time $\geq t$
 - once a node is scheduled, delete it and update the next node in the graph with
$$\text{MST} = \text{wt}_{\text{node}} + \text{start time}_{\text{node}}$$
4. Repeat till all nodes are scheduled (i.e. DAG is empty)

* Due to the MST update step, we ensure that solution given by greedy is feasible.

To prove: $T \leq 2 \text{OPT}$

Proof: Let d^* be the wt of the max lot path in the DAG.

e.g.  $d^* = 4$

then $\text{OPT} \geq d^*$ — ① lower bound for OPT

and $\text{OPT} \geq n/m$ — ② lower bound for OPT

Let's assume $T > 2 \text{OPT}$ and consider the first node j^* due to which the makespan exceeded 2OPT .

collecti of jobs in DAG

Since each job has processing time $P_j = 1$

$\Rightarrow \geq 2 \text{OPT}$ jobs are scheduled before j^*

Consider 2 cases :-

(i) Based on the DAG, the MST of $j^* = 2 \text{OPT}$

i.e. the chain in the DAG has 2OPT jobs before it (this is irrespective of m)

which means that $d^* > \text{MST of } j^*$

$$\Rightarrow d^* > 2 \text{OPT}$$

Not possible, violates ①

(ii) Based on DAG, $\text{MST of } j^* < 2 \text{OPT}$

which means that \exists a node i^* ($\text{MST } i^* \leq \text{MST } j^*$) s.t. it was scheduled later than its MST in the DAG.

\rightarrow because all m machines were not available

\rightarrow At this stage List Scheduling is used. And we know that LS is 2 approx.

So it is not possible for $T > 2 \text{OPT}$

Problem 3 (10 points) In the K-CENTER problem, we are given a set of n points P . We are also given a distance function $d(u, v) \geq 0$ between every pair of points $u, v \in P$. Further, the distance function is a *metric*, meaning they satisfy the following properties :

- (reflexive) $d(u, u) = 0$
- (symmetric) $d(u, v) = d(v, u)$
- (triangle inequality) $d(u, v) \leq d(u, x) + d(x, v)$ for all $x \in P$

The third property is very crucial to the problem (and so is the second !). Now, the goal is to determine a set of k centers $C \subseteq P$. Each point in P will now assign themselves to the closest selected center to form a *clustering*. The goal is to select C in such a way that minimizes the maximum distance from any point to its cluster center.

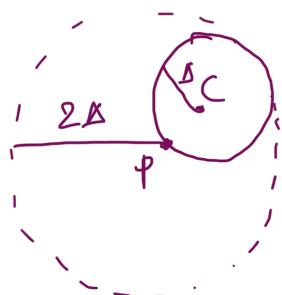
- Suppose we know the value of $\text{OPT} = \Delta$. Here is an algorithm to solve the problem. Pick any arbitrary point, say v as the first center. Assign every point within distance of 2Δ from v to the first center v (including v itself) and delete them. Continue the same procedure on the remaining instance until all points are either a center or assigned to some center.
Prove the given Δ is the correct value of OPT , the above algorithm does not need more than k iterations.
- How does the above imply a polynomial time algorithm to solve the actual problem (where you do not know Δ) ?

a) First, we will try to prove why the algo works.

Let K = be the set of centers chosen by OPT

For any arbitrary point p , the distance for some $c \in K$
 is $d(c, p) \leq \Delta$

So if we choose p as a centre in the algorithm, c and all the other points covered by c in OPT will be covered by p .



$$\begin{aligned} &\because \text{if for some pt } v \text{ s.t.} \\ & \quad d(c, v) \leq \Delta \text{ then} \\ & \quad d(v, p) \leq d(v, c) + d(c, p) \\ & \quad \leq 2\Delta \end{aligned}$$

thus each arbitrary selection will atleast cover one of the centers c (which was included in OPT).

there can be k such centers, so the algorithm takes almost k iterations (here its crucial that covered points be deleted)

* Also its not possible that the center of an arbitrary point (c in OPT) is not covered when the arbitrary point is chosen.

b) K-CENTER(k, P):

$$LB = \min_{v \in P} d(v, v) \quad \forall v \in P$$

$$UB = \max_{v \in P} d(v, v) \quad \forall v \in P$$

while $LB < UB$:

$$\text{set } \Delta \leftarrow \frac{LB + UB}{2}$$

\rightarrow decides based on if algo in (a) ends after k iterations

$\text{isFeasible} = \text{KCENTER-Dec}(P, k, \Delta)$

if $\text{isFeasible} = \text{True}$:

$$UB = \Delta$$

else

$$LB = \Delta$$

Problem 4 (10 points) Suppose we are given a *directed acyclic graph* with edge costs c_e . Further, there is a length ℓ_e associated with every edge - assume all c_e, ℓ_e are non-negative. Given two vertices s and t , we need to find the $s - t$ path of minimum *cost* under the constraint that the *length* of the path does not exceed a given budget L . Design a polynomial time approximation scheme for this problem.

(Will design a PTAS using the same concepts as in KP)

General DP:

$$\begin{aligned} \text{OPT}(v, l) &= \min \text{ cost path from } s \text{ to } v \text{ of length at most } l \\ &= \min \left\{ \begin{array}{l} \text{OPT}(v, l-1) \\ \min_{u \in V} \left\{ \text{OPT}(u, l - \ell_{uv}) + c_{uv} : l - \ell_{uv} \geq 0 \right. \end{array} \right. \end{aligned}$$

$$\text{Answer} = \text{OPT}(t, L)$$

this gives us a time complexity of $O(n^2 L)$

which is pseudopolynomial due to dependence on L which can be arbitrarily large.

Let's define a different DP based on cost

$$\begin{aligned} \text{LC}(v, c) &= \min \text{ length of a path from } s \text{ to } v \text{ with cost atmost } c \\ &= \min \left\{ \begin{array}{l} \text{LC}(v, c-1) \\ \min_{u \in V} \left\{ \text{LC}(u, c - c_{uv}) + \ell_{uv}, \quad c - c_{uv} \geq 0 \right. \end{array} \right. \\ &= \infty, \quad \text{if } c = 0 \text{ and } v \neq s \quad \text{--- (1)} \\ &= 0, \quad v = s \end{aligned}$$

Time complexity: $O(n^2 C) \Rightarrow O(n^3 C_{\max})$ $C = \sum_{e \in E} c_e$

Space complexity : $O(nC)$

Answer: c s.t. $LC(t, c) = \min_{1 \leq c' \leq C} LC(t, c')$

[if $LC(t, c) > L$, then there is no feasible solution]

Scaling the instance

Scale each c_e s.t. $\hat{c}_e \in [1, n/\epsilon]$

Scaling factor, $k = \frac{C_{\max} \epsilon}{n}$ ($n = |V|$)

PTAS algorithm:

1. Choose $\epsilon > 0$

2. Scale $\hat{c}_e = \left\lceil \frac{c_e}{k} \right\rceil$

3. Solve scaled instance using ①

Approximation

$$\begin{aligned}
 \text{cost}(LC) &\leq k \left[\sum_{e \in O} \hat{c}_e \right] \xrightarrow{\text{scaled instance}} \\
 &\leq k \text{ cost(opt)} \\
 &= k \sum_{e \in O} \hat{c}_e \\
 &= k \sum_{e \in O} \left\lceil \frac{c_e}{k} \right\rceil \\
 &\leq \sum_{e \in O} c_e + n k \quad \begin{array}{l} \text{shortest path will have} \\ \text{at most } n \text{ (actually } n-1) \\ \text{edges} \end{array} \\
 &\leq OPT + \epsilon C_{\max} \\
 &\leq OPT(1 + \epsilon)
 \end{aligned}$$

Runtime using scaled instances $O(n^3 \times n/\epsilon) \sim O(n^4/\epsilon)$